

Canny Edge Detector

Algorithms Used:

- I. For Gaussian Smoothing:
 - A. Load training image and convert it to a grayscale floating-point image.
 - B. A Gaussian kernel is a matrix where the values are computed based on the Gaussian distribution. It gives higher weights to pixels near the center and lower weights to those farther away, creating a "blurring" effect.
 - C. Use the `gaussian_smoothing()` function to convolve the image with a Gaussian kernel.
- II. For Calculating Image Gradient:
 - A. Applying Sobel operators to the smoothed image: Sobel operator detects changes in intensity in the horizontal (G_x) and vertical (G_y) directions and gives I_x and I_y .
 - B. Solving the gradient magnitude (Mag) using the formula: $Mag = \sqrt{I_x^2 + I_y^2}$.
 - C. Solving the gradient direction (Θ) using the formula: $\Theta = \arctan2(I_y, I_x)$.
- III. For Selecting high and low thresholds [`find_threshold()`]:
 - A. Flatten the gradient magnitude matrix into a 1D array.
 - B. Sort the gradient magnitudes in ascending order.
 - C. Set the high threshold (T_{high}) as the value at the `percentageOfNonEdge` percentile.
 - D. So if the percentage of non-edge pixels is 0.7, then we take the 70th percent-th value of the gradient magnitude pixel values as the high threshold as the Mag value represents much of that pixel is likely to be an edge.
 - E. Set the low threshold (T_{low}) as half of the high threshold (T_{high}).
- IV. For Suppressing Nonmaxima [`nonmaxima_suppress()`]:
 - A. Iterate through each pixel in the gradient magnitude matrix (Mag).
 - B. For each pixel, check its two neighbors along the gradient direction (Θ).
 - C. Suppress the pixel if its magnitude is not greater than or equal to both neighbors.
 - D. The result is a thinned edge map where only the strongest edge pixels remain.
- V. Thresholding and Edge Linking [`edge_linking()`]:
 - A. Classify pixels based on the thresholds calculated previously.
 - B. Mark pixels with magnitude greater than or equal to T_{high} as strong edges.
 - C. Mark pixels with magnitude between T_{low} and T_{high} as weak edges.
 - D. Set all other pixels to zero.
 - E. For each weak edge pixel, check its 8-connected neighbors.
 - F. Promote the weak edge to a strong edge if any neighbor is a strong edge.
 - G. Suppress the weak edge if no strong edge is found in its neighbors.
 - H. The final edge map contains only strong edges (255) and non-edges (0).


- VI. Comparisons of different parameters
- VII. Comparisons with different edge detectors.

Results:

1. Input Test Images:




2. With Gaussian filtering:

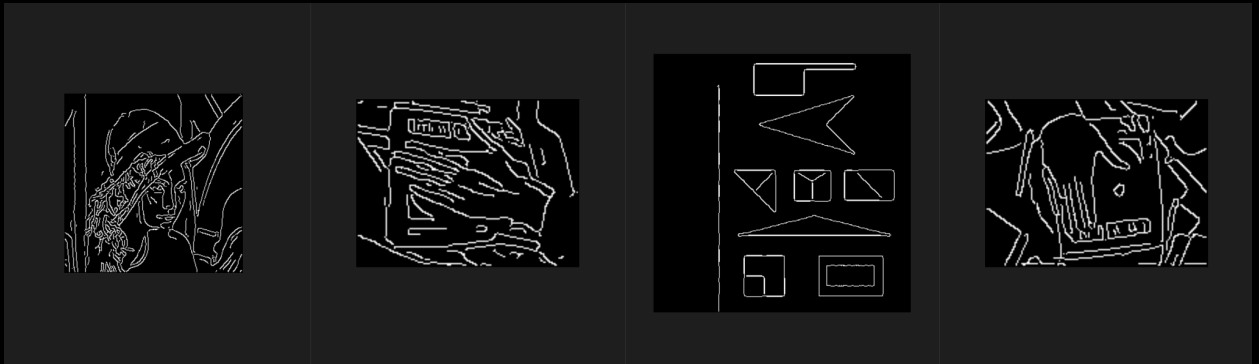


3. With (N, Sigma) as (5, 1.4) and percentageOfNonEdge = 0.7

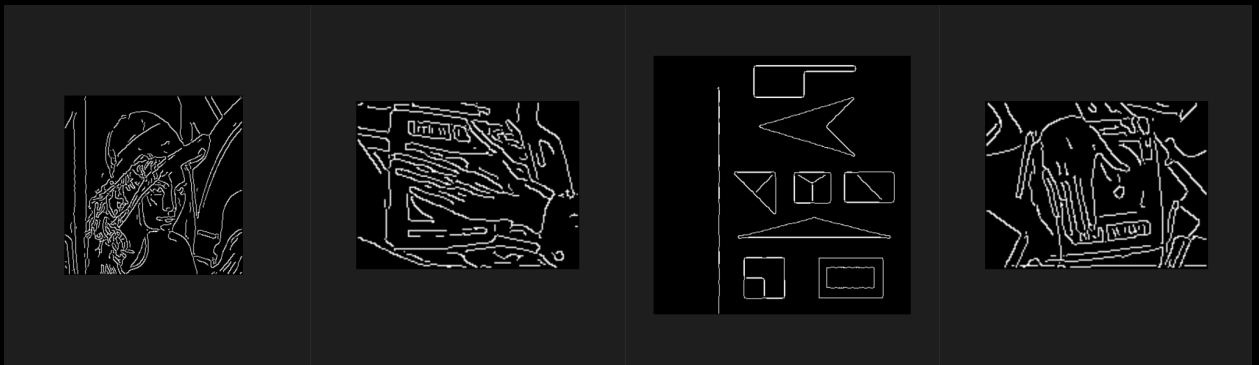
Based on the observation on test.bmp it seems that Over-smoothing: The output seems to have lost nearly all edge information, suggesting the edges were smoothed away too aggressively before gradient computation. And Thresholding was too aggressive: It's likely the T_{high} threshold is too high due to a high percentageOfNonEdge(0.7), causing valid weak edges to be ignored. **And the main issue is** that the gradient of most pixel is not too high as neighbouring colours are pretty similar, so I added a min-max thresholding which gave much better results.



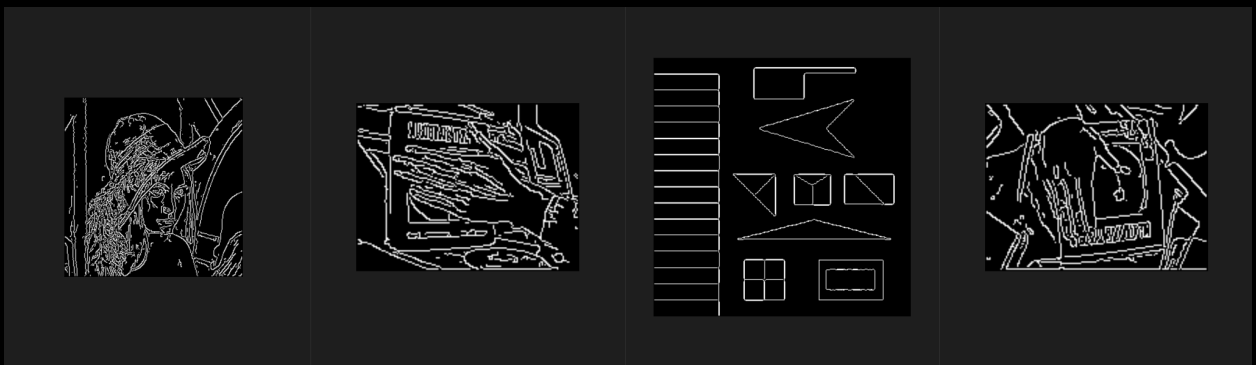
4. With (N, Sigma) as $(5, 1.4)$ and $\text{percentageOfNonEdge} = 0.7$ and thresholding.



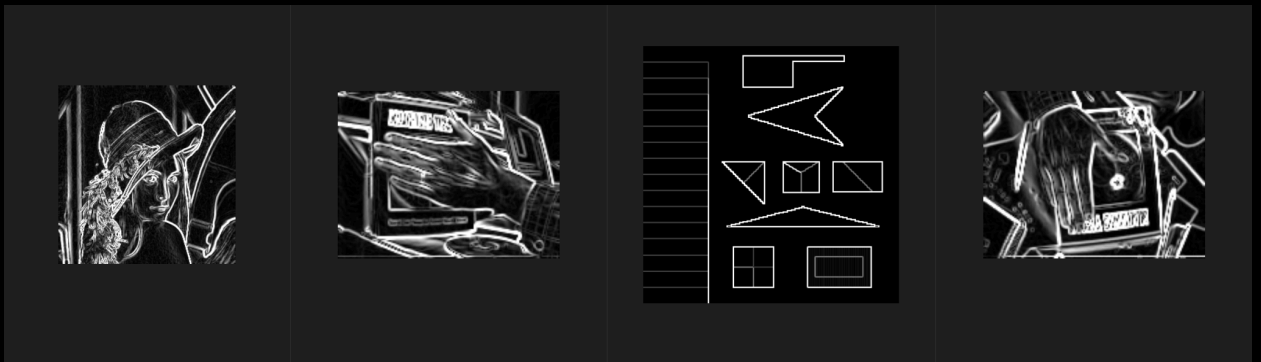
5. With (N, Sigma) as $(5, 1.4)$ and Lowering $\text{percentageOfNonEdge}$ to 0.2. This will lower T_{high} and as we can observe retains more edge options.



6. With (N, Sigma) as $(3, 0.4)$ and $\text{percentageOfNonEdge} = 0.7$. This gives a much noisier output.



7. Comparison with CV's Sobel function:



8. Comparison with CV's Canny function:

